

Introduction

— Boston University CS 506 - Lance Galletti —



Data Science

- Collection of methods and tools that allow for extracting knowledge from data
- Cross-disciplinary:
 - Math
 - Statistics
 - Computer Science
 - Domain Expertise
- Know what you don't know!

Knowledge = Testable Predictions



Knowledge = Testable Predictions



at time t_0

Knowledge = Testable Predictions

Model:

$f(x, y, t) \Rightarrow \text{temperature}$

Magic


VS

“Heat Diffusion”

Which theory should we use?
How to distinguish or unify them?

Knowledge = Testable Predictions

Scientific perspective: look at
what each theory anticipates!



Knowledge = Testable Predictions

If you can equally well explain every outcome, how can you have a definitive / deterministic anticipation of events?

If you're equally good at explaining every outcome, you have zero knowledge.

Confirmation Bias

In a class just like this one, imagine playing the following game...

Confirmation Bias

I announce “(2, 4, 6) follows the rule”.

Here are the examples submitted by one of the participants:

- (2, 4, 3) -> NO
- (6, 8, 10) -> YES
- (1, 3, 5) -> YES

After which, they proceed to write down their hypothesized rule. Would you have wanted to try more examples? If so, which and for what reason?

Confirmation Bias

Let's take a poll:

- A. (100, 102, 104)
- B. (5, 7, 9)
- C. (1, 2, 3)

Confirmation Bias

Challenges of Data Science:

- Not all examples contribute similar amounts of information
- A set of examples may not always be representative of the underlying rule
- There may be infinitely many rules that match the examples provided

So Data Science is VERY DIFFICULT!!! All models are wrong but some are useful

Confirmation Bias

Positive Examples VS Negative Examples

assuming the hypothesis h is $(x, x+2, x+4)$ which type of examples are the following:


- $(2, 4, 3)$
- $(6, 8, 10)$
- $(1, 3, 5)$

Confirmation Bias

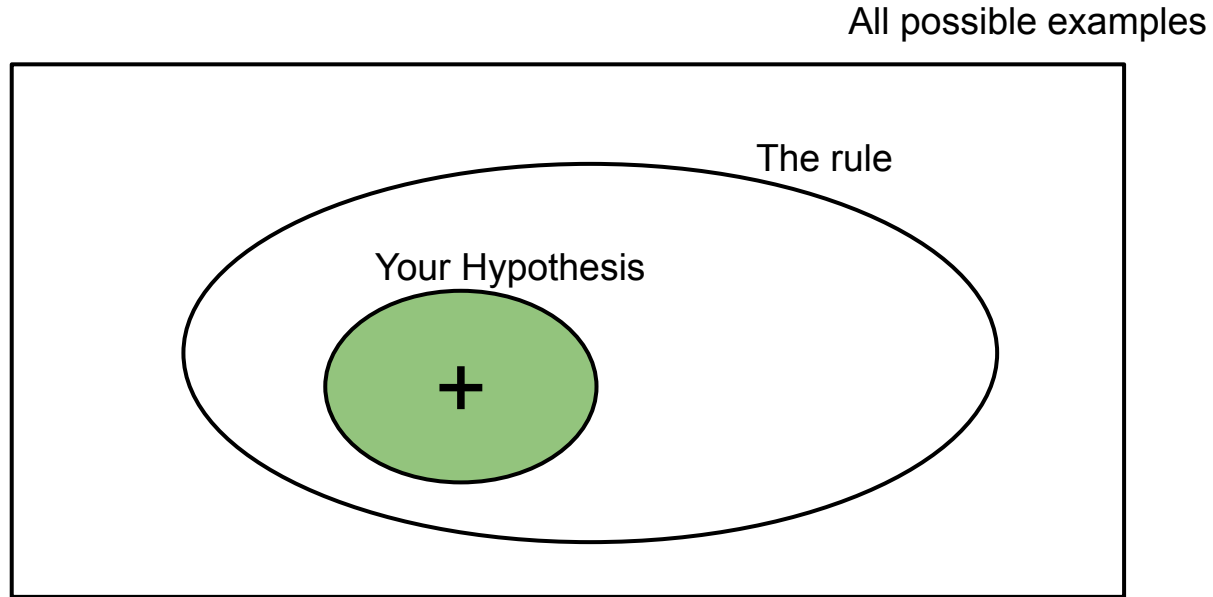
- Both positive and negative examples can falsify a hypothesis
- Tendency to choose positive ones over negative ones

Confirmation Bias


All possible examples




Confirmation Bias



Confirmation Bias



Confirmation Bias



Confirmation Bias

Let's take a poll:

- A. (100, 102, 104)
- B. (5, 7, 9)
- C. (1, 2, 3)

Confirmation Bias

The rule was ($a < b < c$).

If you only tried positive examples of either (x, x + 2, x+4) or (x, 2x 3x) you would only get confirmation.

For reference, this exercise was first introduced by Wason P.C in 1960 as part of a journal in experimental psychology.

Confirmation Bias



Types of Data

Types of Data - Records


m-dimensional points / vectors

Example: (name, age, balance) -> ("John", 20, 100)

Types of Data - Records

m-dimensional points / vectors


Example: (name, age, balance) \rightarrow ("John", 20, 100)



Types of Data - Graphs


Nodes connected by edges

Example:



Adjacency Matrix


	1	2	3
1	0	1	1
2	1	0	1
3	1	1	0





Adjacency List

1 : {2, 3}
2 : {1, 3}
3 : {1, 2}

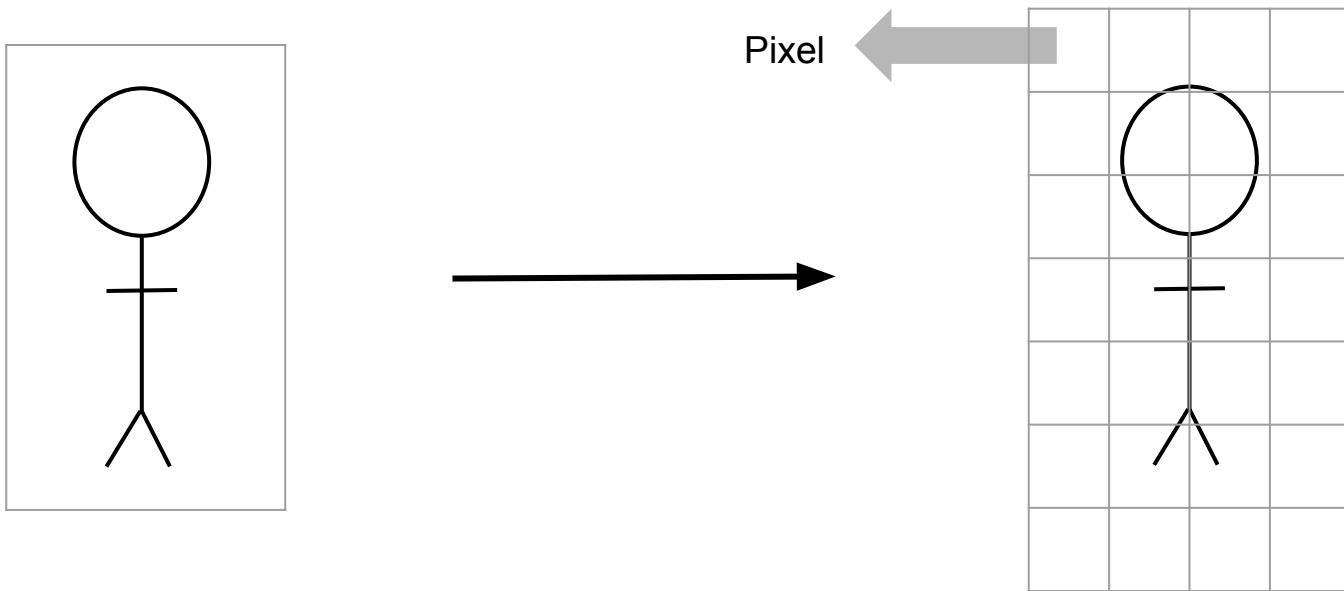
Types of Data - Images




Types of Data - Images



Types of Data - Images



Types of Data - Text



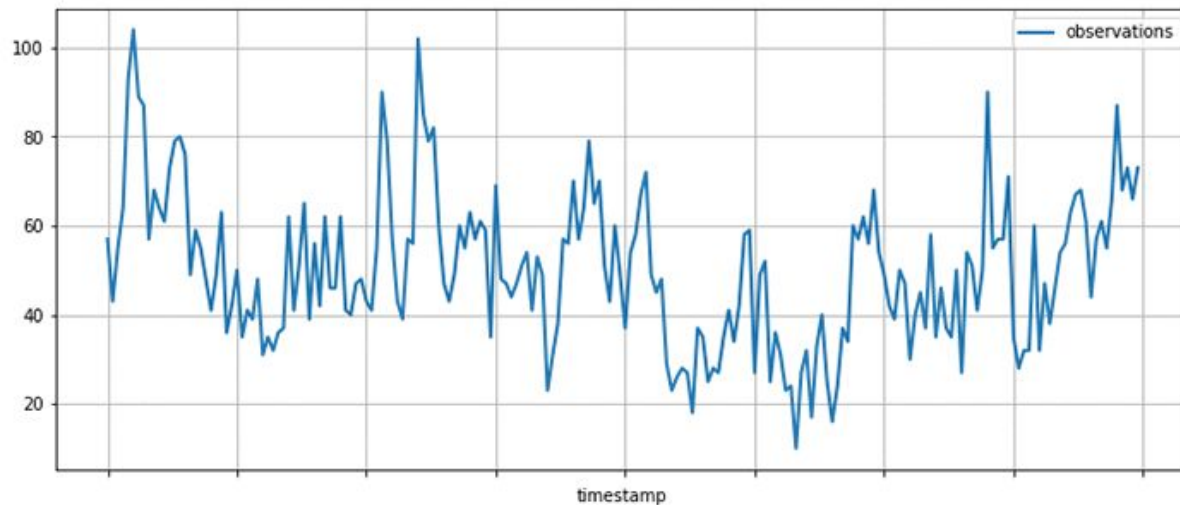
List of words

Types of Data - Strings

DNA seq (A T G C C G T A ...) -> list of characters

Types of Data - Time Series

List of data at specific intervals of time




Types of Learning

- Unsupervised Learning
- Supervised Learning


Unsupervised Learning

Goal: Find interesting structure in the data



Unsupervised Learning

Goal: Find interesting structure in the data



This type of unsupervised learning is referred to as clustering

Unsupervised Learning

Dataset: Collection of Articles

Question: Are these articles covering the same topics?

Unsupervised Learning

Goals:


1. Better understand / describe the data
 - a. Data exploration / visualization step
 - b. Recommender Systems (similar users might be recommended the same things, emails similar to those marked as spam could be spam etc.)
2. Provide sensible defaults to missing values
 - a. Data preprocessing step

Supervised Learning

cricket chirps / min	temperature
10	40
5	37
17	53
55	103
40	78

Supervised Learning

cricket chirps / min	temperature
10	40
5	37
17	53
55	103
40	78




Supervised Learning

cricket chirps / min	temperature
10	40
5	37
17	53
55	103
40	78




Supervised Learning

cricket chirps / min	temperature
10	40
5	37
17	53
55	103
40	78




Supervised Learning

cricket chirps / min	temperature
10	40
5	37
17	53
55	103
40	78




Supervised Learning

cricket chirps / min	temperature
10	40
5	37
17	53
55	103
40	78



Supervised Learning

cricket chirps / min	temperature
10	40
5	37
17	53
55	103
40	78




This type of supervised learning is referred to as regression

Supervised Learning

age	tumor size	malignant
20	12	0
22	15	1
47	20	1
59	2	1


Supervised Learning

age	tumor size	malignant
20	12	0
22	15	1
47	20	1
59	2	1




Supervised Learning

age	tumor size	malignant
20	12	0
22	15	1
47	20	1
59	2	1




Supervised Learning

age	tumor size	malignant
20	12	0
22	15	1
47	20	1
59	2	1




This type of supervised learning is referred to as classification

Data Science Workflow (simplified)



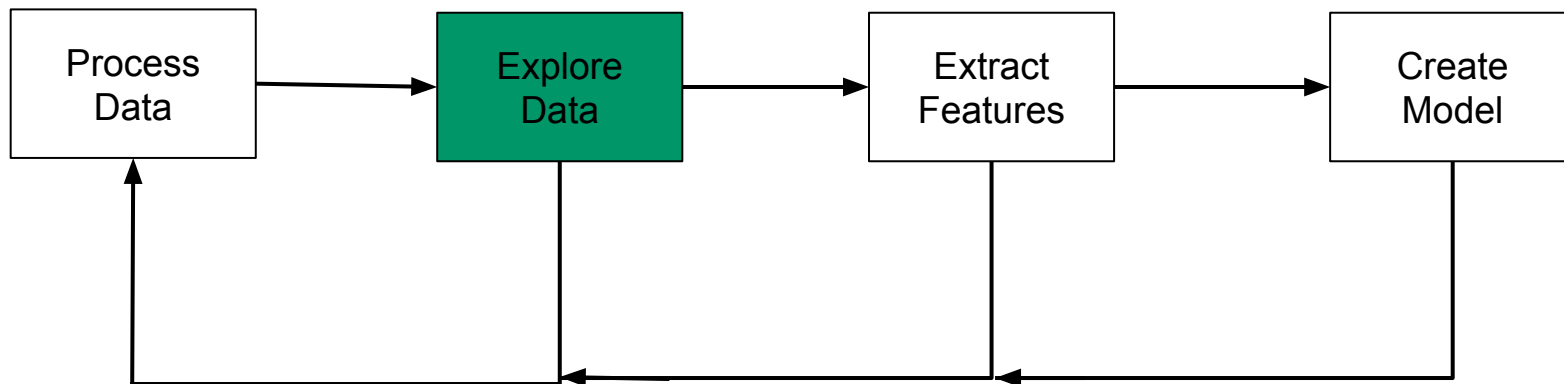
Data Science Workflow (simplified)



Data Processing

- What data should and shouldn't be used for the task?
- What to do with missing data?
- What to do with inconsistent data?
- What assumptions are you making with the transformations of the data?


Data Science Workflow (simplified)



Exploratory Data Analysis

- Describe, contextualize, and visualize the data
- What might be related to what you're trying to predict?
- Are there imbalances in the data?


Data Science Workflow (simplified)



Feature Extraction

- Are the features provided by the dataset the best features to use for the task?
- What other features can be extracted?
- Should existing features be transformed?

Data Science Workflow (simplified)



Finding the right model

- Ask what and **who** the model is used / intended for
 - Is it just the general trend that is important or the exact predictions that are important?
 - Is this a problem that **needs** predictive tools to solve?
- The success of this step depends entirely on the work done in previous steps
 - remember: **garbage in, garbage out!** (it's all about the data)
- Is your model easy to explain?
- When your model fails, can you explain why?